

## Lab Work 5: Time Series Forecasting with LSTM

**Name:** Al Amin Hossain Nayem

**Date:** April 6, 2025

---

### Objective

The goal of this lab work is to implement a Long Short-Term Memory (LSTM) model to forecast SO<sub>2</sub> (Sulfur Dioxide) concentrations 1 to 4 hours ahead based on past observations. The model is trained on January 2018 data and tested on February 2018 data.

### Tools and Libraries

- Python
- PyTorch
- NumPy
- Matplotlib

### Dataset

- **Training set:** 2018\_01\_so2.npy
- **Testing set:** 2018\_02\_so2.npy

### Code Overview

#### Dataset Preparation

The TimeSeriesDataset class is used to generate sequences:

- Input window size: 24 hours
- Output window size: 4 hours

#### Model Definition

An LSTM network with:

- Input size = 1
- Hidden size = 128
- 2 layers
- Dropout = 0.2

- Fully connected (FC) layer outputs 4 values for next 4 hours.

## **Normalization**

The data is normalized using the training set's mean and standard deviation.

## **Training**

- Optimizer: Adam
- Learning rate: 0.001
- Loss function: Mean Squared Error (MSE)
- Epochs: 150
- Batch size: 64

## **Testing and Evaluation**

- Total test loss is computed.
- Individual loss is computed for each of the 4 prediction hours.
- True and predicted values are plotted for each forecast horizon.

## **Results**

### **Loss during Training**

Loss steadily decreased across 150 epochs, indicating successful model convergence.

### **Testing Results**

- **Total Test Loss (MSE):** (Shown in the console)
- **Individual Losses:**
  - 1 hour ahead: (Shown in the console)
  - 2 hours ahead: (Shown in the console)
  - 3 hours ahead: (Shown in the console)
  - 4 hours ahead: (Shown in the console)

## **Plots**

Graphs were generated comparing true vs. predicted SO<sub>2</sub> concentrations for:

- 1 hour ahead

- 2 hours ahead
- 3 hours ahead
- 4 hours ahead

These plots demonstrate that the model captures the temporal pattern well but with increasing error as prediction horizon extends.

## Conclusion

An LSTM-based model was successfully built and trained to forecast short-term SO<sub>2</sub> concentrations. The model shows good predictive performance for immediate future steps but degrades slightly for farther steps. The performance could be further improved by:

- Tuning hyperparameters (e.g., hidden size, learning rate)
- Using more complex architectures (e.g., attention mechanisms)
- Training on more data.

## Figures:

# -\*- coding: utf-8 -\*-

"""

Created on Sun Apr 6 10:49:57 2025

@author: ginta

"""

import torch

import torch.nn as nn

import torch.optim as optim

from torch.utils.data import Dataset, DataLoader

import numpy as np

import matplotlib.pyplot as plt

# Define the Dataset

```
class TimeSeriesDataset(Dataset):
```

```
    def __init__(self, series, input_window=24, output_window=4):
```

```
        self.series = series
```

```
        self.input_window = input_window
```

```
        self.output_window = output_window
```

```
        self.samples = []
```

```
        for i in range(len(series) - input_window - output_window + 1):
```

```
            x = series[i : i + input_window]
```

```
            y = series[i + input_window : i + input_window + output_window]
```

```
            self.samples.append((x, y))
```

```
    def __len__(self):
```

```
        return len(self.samples)
```

```
    def __getitem__(self, index):
```

```
        x, y = self.samples[index]
```

```
        return torch.tensor(x, dtype=torch.float32).unsqueeze(-1), torch.tensor(y,  
dtype=torch.float32)
```

```
# Define the Model
```

```
class LSTMModel(nn.Module):
```

```
    def __init__(self, input_size=1, hidden_size=128, num_layers=2, output_size=4):
```

```
        super(LSTMModel, self).__init__()
```

```
        self.hidden_size = hidden_size
```

```
        self.num_layers = num_layers
```

```
self.lstm = nn.LSTM(input_size, hidden_size, num_layers, batch_first=True,
dropout=0.2)
```

```
self.fc = nn.Linear(hidden_size, output_size)
```

```
def forward(self, x):
```

```
    h0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(x.device)
```

```
    c0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(x.device)
```

```
    out, _ = self.lstm(x, (h0, c0))
```

```
    out = out[:, -1, :]
```

```
    out = self.fc(out)
```

```
    return out
```

```
# Load and normalize the data
```

```
train = np.load('2018_01_so2.npy')
```

```
test = np.load('2018_02_so2.npy')
```

```
# Normalize
```

```
train_mean = train.mean()
```

```
train_std = train.std()
```

```
train = (train - train_mean) / train_std
```

```
test = (test - train_mean) / train_std
```

```
# Create Dataset and DataLoader
```

```
dataset = TimeSeriesDataset(train, input_window=24, output_window=4)
```

```
dataloader = DataLoader(dataset, batch_size=64, shuffle=True)
```

**# Hyperparameters**

**num\_epochs = 150**

**learning\_rate = 0.001**

**# Model, Loss, Optimizer**

**model = LSTMModel(input\_size=1, hidden\_size=128, num\_layers=2, output\_size=4)**

**criterion = nn.MSELoss()**

**optimizer = optim.Adam(model.parameters(), lr=learning\_rate)**

**# Training loop**

**for epoch in range(num\_epochs):**

**model.train()**

**epoch\_loss = 0.0**

**for x\_batch, y\_batch in dataloader:**

**optimizer.zero\_grad()**

**predictions = model(x\_batch)**

**loss = criterion(predictions, y\_batch)**

**loss.backward()**

**optimizer.step()**

**epoch\_loss += loss.item()**

**avg\_loss = epoch\_loss / len(dataloader)**

**print(f"Epoch {epoch+1}/{num\_epochs}, Loss: {avg\_loss:.6f}")**

**# Testing**

```

model.eval()

dataset_test = TimeSeriesDataset(test, input_window=24, output_window=4)

full_batch_loader = DataLoader(dataset_test, batch_size=len(dataset_test),
                                shuffle=False)

for x_full, y_full in full_batch_loader:

    with torch.no_grad():

        predictions = model(x_full)

# Calculate total test loss

total_loss = criterion(predictions, y_full)

print(f"\nTotal Test Loss (MSE): {total_loss:.6f}")

# Calculate individual losses for each of the 4 output steps

for i in range(4):

    step_loss = criterion(predictions[:,i], y_full[:,i])

    print(f"Test Loss for {i+1} hour(s) ahead: {step_loss:.6f}")

# Convert predictions and targets back to numpy (and denormalize for plotting if
needed)

Predictions = predictions.cpu().numpy() * train_std + train_mean

Targets = y_full.cpu().numpy() * train_std + train_mean

# Plot true vs predicted for each forecast step

for i in range(4):

    plt.figure(figsize=(10, 4))

    plt.plot(Targets[:,i], label="True Values")

```

```
plt.plot(Predictions[:,i], label="Predicted Values")  
plt.title(f"SO2 Concentration Prediction {i+1} hour(s) ahead")  
plt.xlabel("Sample Index")  
plt.ylabel("SO2 Concentration")  
plt.legend()  
plt.grid(True)  
plt.show()
```



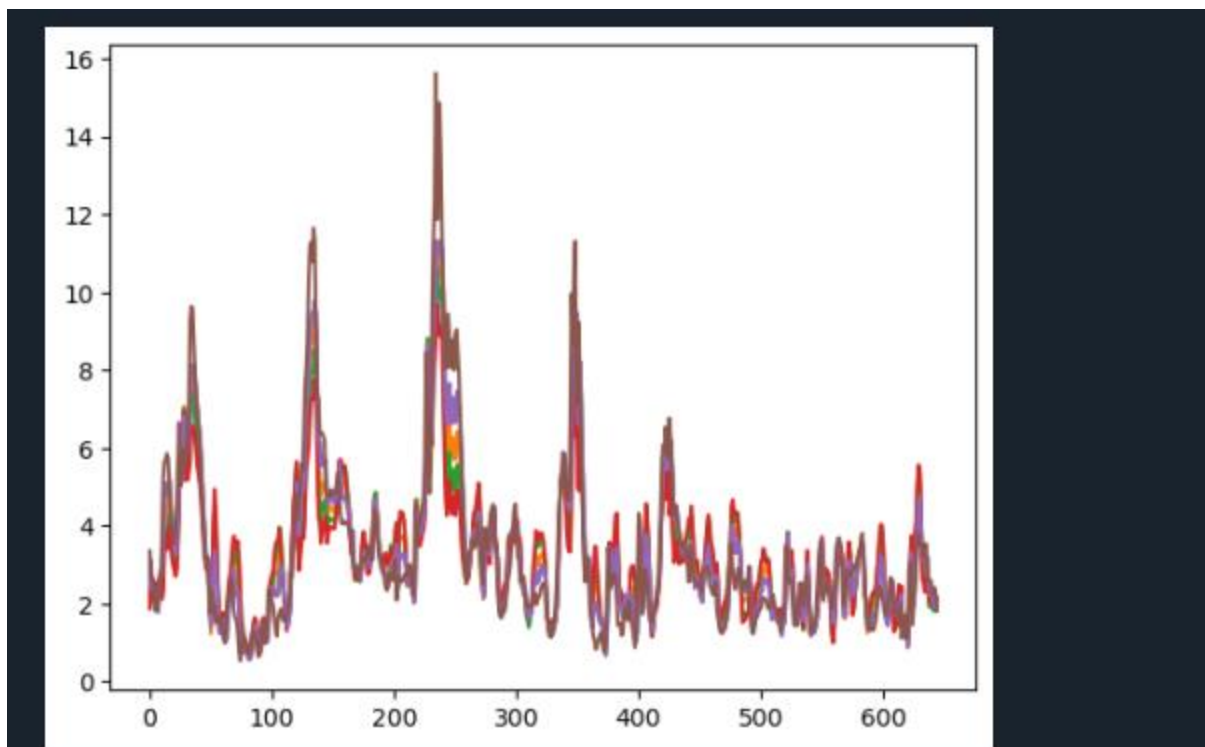
Console 1/A X

Epoch 142/150, Loss: 0.017895  
Epoch 143/150, Loss: 0.019308  
Epoch 144/150, Loss: 0.021125  
Epoch 145/150, Loss: 0.021733  
Epoch 146/150, Loss: 0.018767  
Epoch 147/150, Loss: 0.018976  
Epoch 148/150, Loss: 0.017991  
Epoch 149/150, Loss: 0.018577  
Epoch 150/150, Loss: 0.020029  
  
Total Test Loss (MSE): 0.692901  
Test Loss for 1 hour(s) ahead: 0.287682  
Test Loss for 2 hour(s) ahead: 0.575706  
Test Loss for 3 hour(s) ahead: 0.827717  
Test Loss for 4 hour(s) ahead: 1.080501

Name	Type	Size	Value
avg_loss	float	1	0.020028527670850355
dataloader	utils.data.dataloader.DataLoader	12	DataLoader object of torch.utils.data.dataloader module
dataset	TimeSeriesDataset	717	TimeSeriesDataset object of __main__ module
dataset_test	TimeSeriesDataset	645	TimeSeriesDataset object of __main__ module
epoch	int	1	149
epoch_loss	float	1	0.24034233205020428
full_batch_loader	utils.data.dataloader.DataLoader	1	DataLoader object of torch.utils.data.dataloader module
i	int	1	3
learning_rate	float	1	0.001
loss	Tensor	1	Tensor object of torch module
num_epochs	int	1	150
optimizer	optim.adam.Adam	1	Adam object of torch.optim.adam module
predictions	Tensor	(645, 4)	Tensor object of torch module
Predictions	Array of float32	(645, 4)	[[3.3887262 3.4917698 3.2151942 2.752795 ] [2.957333 2.4283543 2.040 ...
step_loss	Tensor	1	Tensor object of torch module
Targets	Array of float32	(645, 4)	[[3.3564901 2.653337 2.6919339 2.5543158] [2.653337 2.6919339 2.554 ...
test	Array of float32	(672,)	[-0.527278 -0.47309926 -0.70915896 ... -0.7410652 -0.7592201 -0.79 ...

Help Variable Explorer Plots Files

Name	Type	Size	Value
learning_rate	float	1	0.001
loss	Tensor	1	Tensor object of torch module
num_epochs	int	1	150
optimizer	optim.adam.Adam	1	Adam object of torch.optim.adam module
predictions	Tensor	(645, 4)	Tensor object of torch module
Predictions	Array of float32	(645, 4)	[[3.3887262 3.4917698 3.2151942 2.752795 ] [2.957333 2.4283543 2.040 ...
step_loss	Tensor	1	Tensor object of torch module
Targets	Array of float32	(645, 4)	[[3.3564901 2.653337 2.6919339 2.5543158] [2.653337 2.6919339 2.554 ...
test	Array of float32	(672,)	[-0.527278 -0.47309926 -0.70915896 ... -0.7410652 -0.7592201 -0.79 ...
total_loss	Tensor	1	Tensor object of torch module
train	Array of float32	(744,)	[-0.35445136 -0.5467527 -0.45554823 ... 0.7760028 -0.25899518 -0.3 ...
train_mean	float32	1	3.0607738
train_std	float32	1	1.8394786
x_batch	Tensor	(13, 24, 1)	Tensor object of torch module
x_full	Tensor	(645, 24, 1)	Tensor object of torch module
y_batch	Tensor	(13, 4)	Tensor object of torch module
y_full	Tensor	(645, 4)	Tensor object of torch module



Spyder (Python 3.12)

File Edit Search Source Run Debug Consoles Projects Tools View Help

D:\AI\labwork\_5\assignment5.py

```
85     epoch_loss += loss.item()
86
87     avg_loss = epoch_loss / len(data_loader)
88     print(f"Epoch {epoch+1}/{num_epochs}, Loss: {avg_loss:.6f}")
89
90 # Testing
91 model.eval()
92 dataset_test = TimeSeriesDataset(test, input_window=24, output_window=4)
93 full_batch_loader = DataLoader(dataset_test, batch_size=len(dataset_test))
94
95 for x_full, y_full in full_batch_loader:
96     with torch.no_grad():
97         predictions = model(x_full)
98
99 # Calculate total test loss
100 total_loss = criterion(predictions, y_full)
101 print(f"\nTotal Test Loss (MSE): {total_loss:.6f}")
102
103 # Calculate individual losses for each of the 4 output steps
104 for i in range(4):
105     step_loss = criterion(predictions[:,i], y_full[:,i])
106     print(f"Test Loss for {i+1} hour(s) ahead: {step_loss:.6f}")
107
108 # Convert predictions and targets back to numpy (and denormalize for plt)
109 Predictions = predictions.cpu().numpy() * train_std + train_mean
110 Targets = y_full.cpu().numpy() * train_std + train_mean
111
112 # Plot true vs predicted for each forecast step
113 for i in range(4):
114     plt.figure(figsize=(10, 4))
115     plt.plot(Targets[:,i], label="True Values")
116     plt.plot(Predictions[:,i], label="Predicted Values")
117     plt.title(f"SD, Concentration Prediction {i+1} hour(s) ahead")
118     plt.xlabel("Sample Index")
119     plt.ylabel("SD, Concentration")
120     plt.legend()
121     plt.grid(True)
122     plt.show()
123
```

16  
14  
12  
10  
8  
6  
4  
2  
0

0 100 200 300 400 500 600

Help Variable Explorer Plots Files

Console 1/A X

```
Epoch 142/150, Loss: 0.017895
Epoch 143/150, Loss: 0.019388
Epoch 144/150, Loss: 0.021125
Epoch 145/150, Loss: 0.021723
Epoch 146/150, Loss: 0.018767
Epoch 147/150, Loss: 0.018976
Epoch 148/150, Loss: 0.017991
Epoch 149/150, Loss: 0.018577
Epoch 150/150, Loss: 0.020029

Total Test Loss (MSE): 0.692901
Test Loss for 1 hour(s) ahead: 0.287682
Test Loss for 2 hour(s) ahead: 0.575706
Test Loss for 3 hour(s) ahead: 0.827717
Test Loss for 4 hour(s) ahead: 1.080501
```

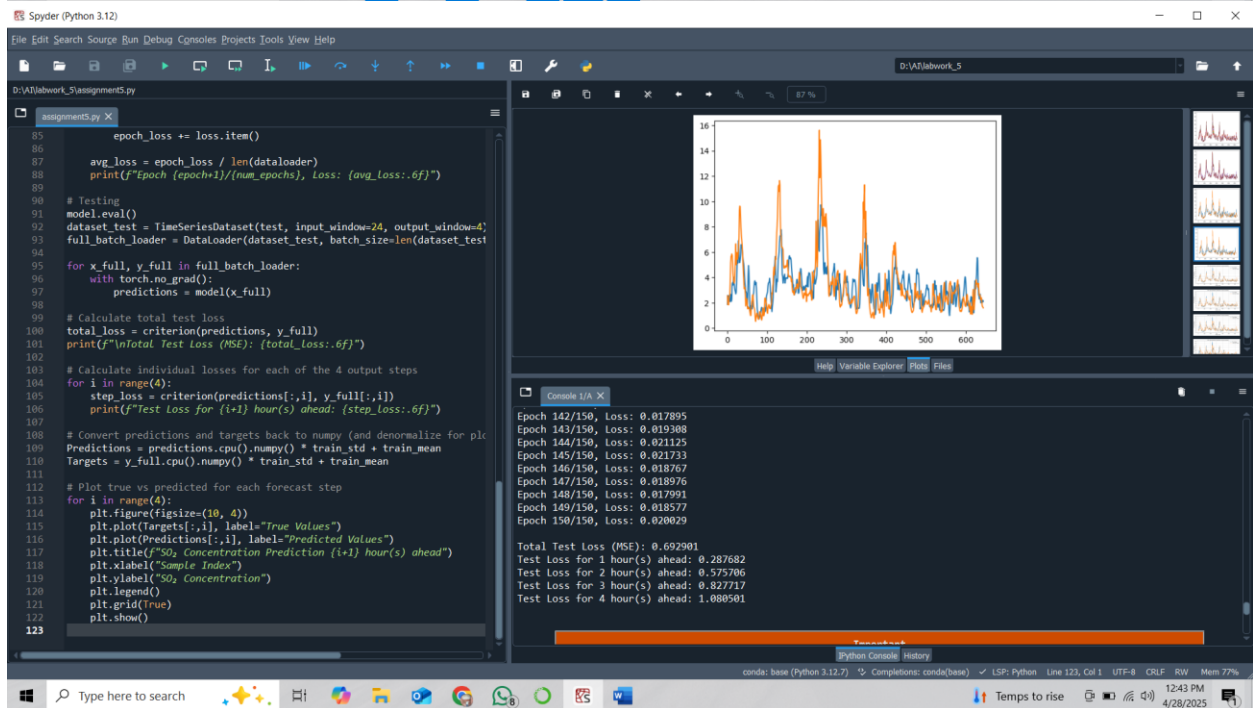
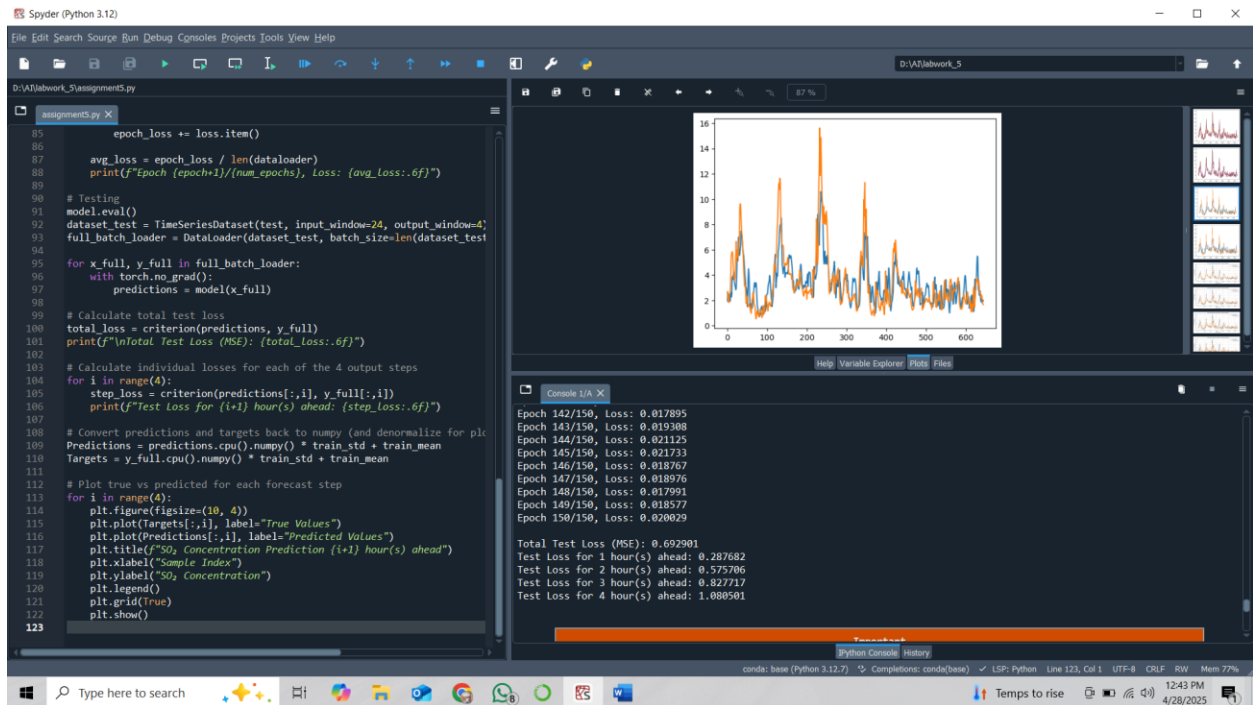
Python Console History

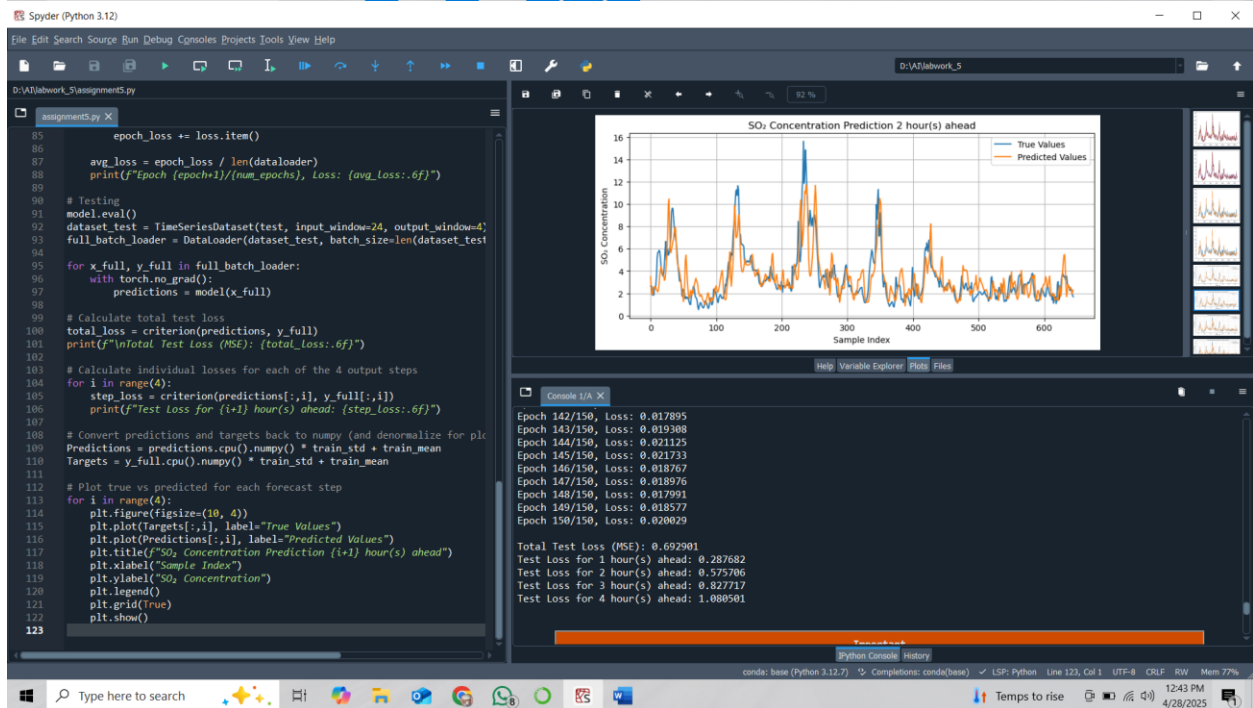
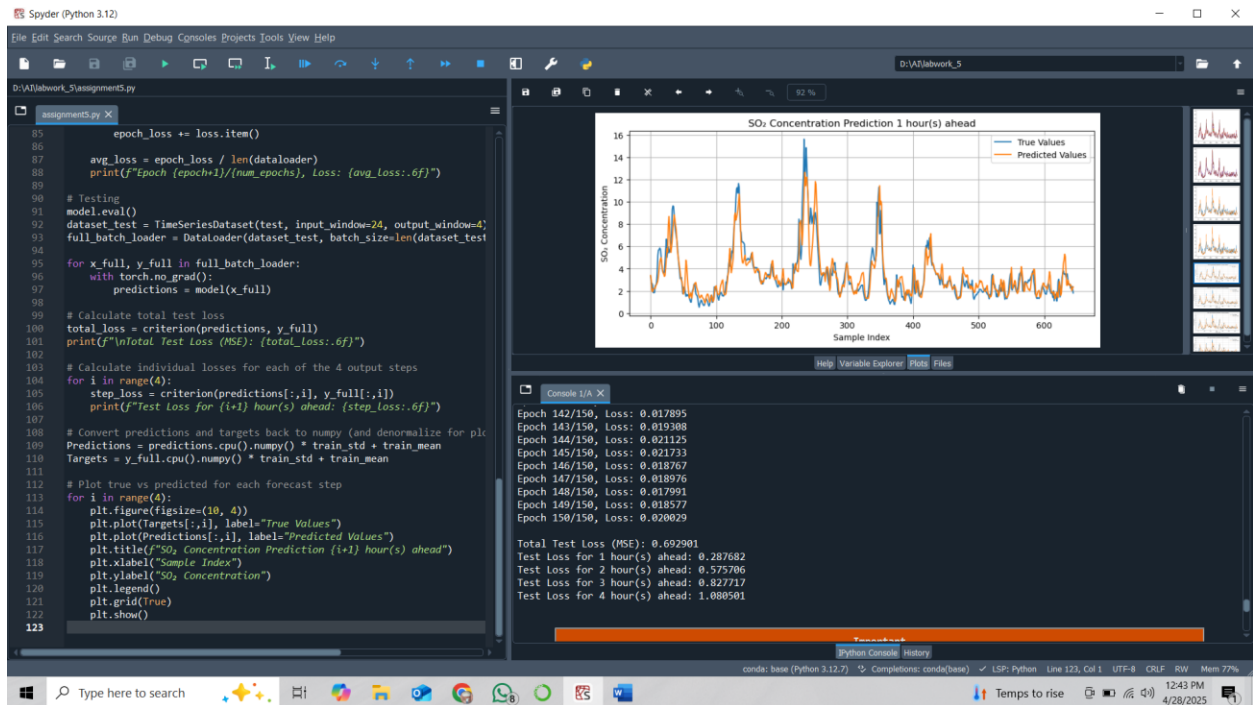
canda: base (Python 3.12.7) ↳ Completions: canda(base) ✓ LSP: Python Line 123, Col 1 UTF-8 CRLF RW Mem 78%

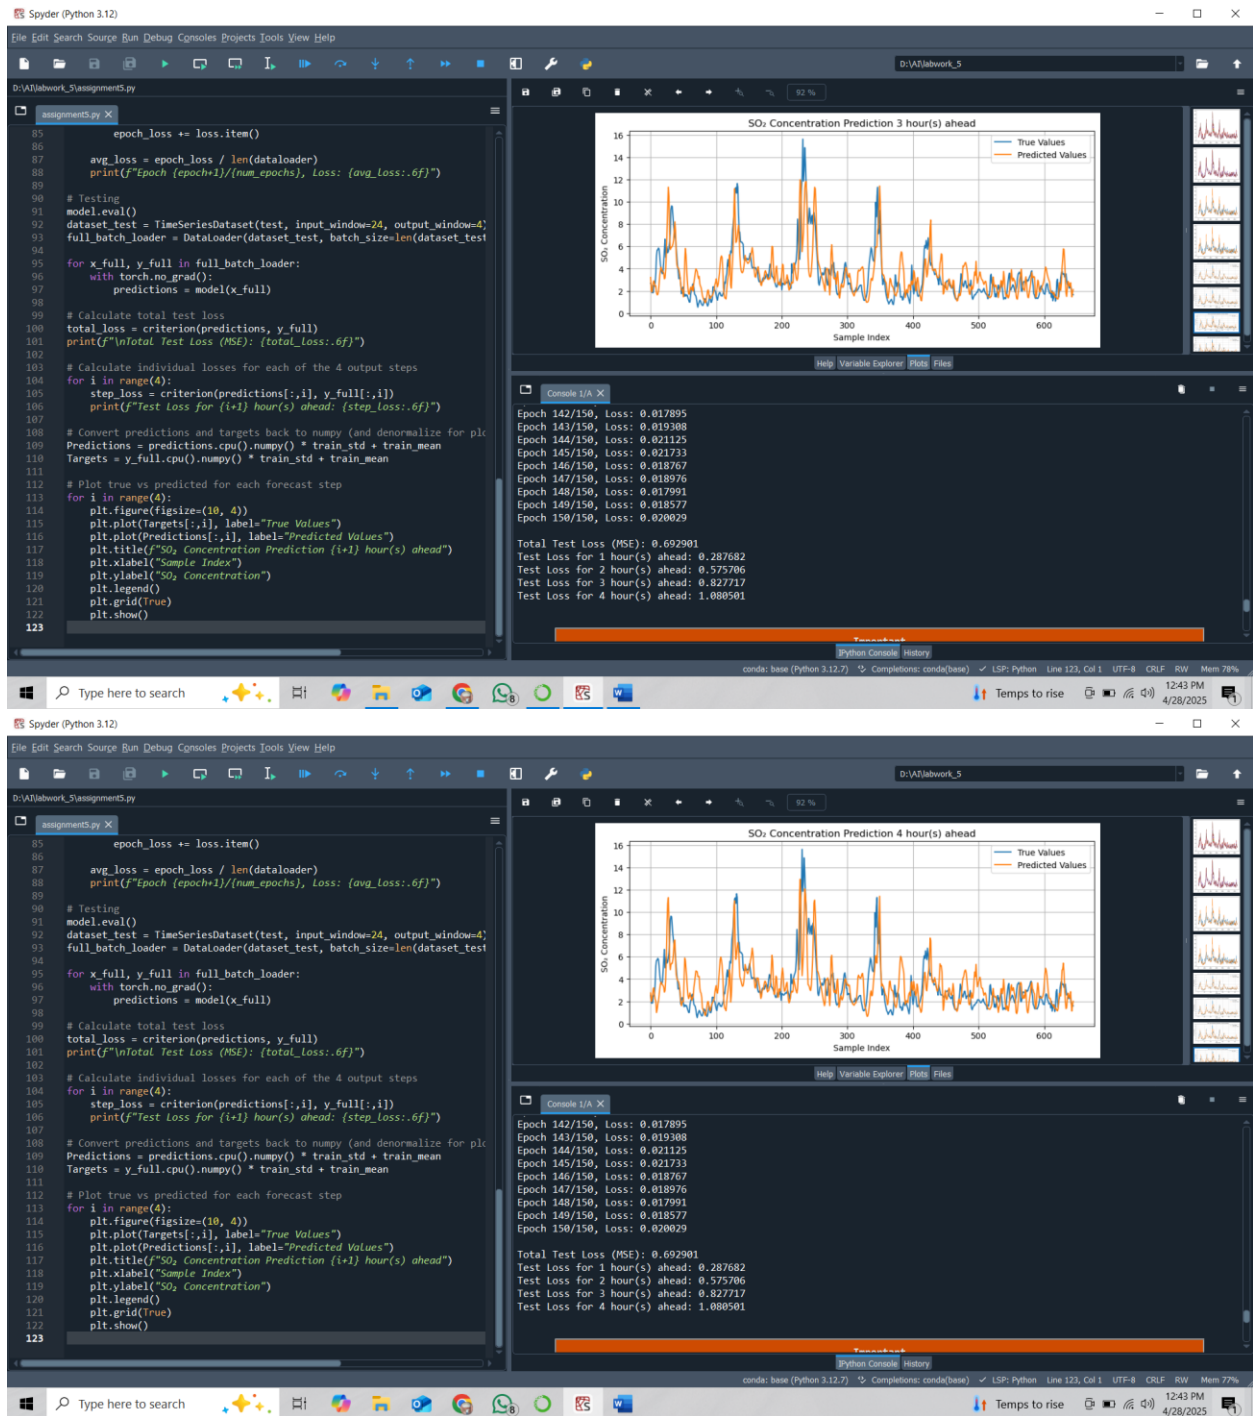
Type here to search

Temps to rise

12:42 PM  
4/28/2025







End of Lab Report